Programming Assignment (Simulating Pipelined Machine)

Phase 1 on-time due: July 2, 2022
Phase 1 late due: July 3, 2022

Phase 2 on-time Due: July 9, 2022
Phase 2 late Due: July 10, 2022

In this assignment, you will become familiar with how a basic MIPS 5-stage pipeline works. You will be given a simulator that models an **unpipelined** processor that implements a small MIPS-like instruction set. **Your assignment is to create a cycle-accurate simulator of a pipelined version of this processor.** Your simulator will perform (1) data forwarding, (2) a simple branch prediction scheme, and (3) the pipeline interlocks to stall the pipeline when necessary.

You will need a Linux environment to complete this assignment. There are multiple ways to set up your work environment. If you use a Mac or a Linux OS, you may already have a working environment or can set up without too much pain (e.g., installing XCode on a Mac or using a package manager like HomeBrew or MacPort). If you use a Windows OS, you can try Windows Subsystem for Linux (link available in the Resource tab in the course homepage) or consider running a Virtual Machine as a guest OS. If you do not want to set up a work environment locally on your computer, you can choose to work on the `grace` cluster as you did before when you took CMSC216. The course space is already created in the cluster and you can login with your directory ID and password (`ssh yourDirectoryID@grace.umd.edu`). If you work on the grace cluster and still want to use Visual Studio Code (VSCode) as your code editor, you can consider installing a package that supports remote SSH connection. It allows you to work on a remote file system in VSCode, like you work on local files.

It is important to understand that there can be subtle differences in the machine configurations in all the options explained above. It is possible that your work run on your work environment but fail on the submit server – you might have encountered such cases when you took CMSC 216. A common cause is the use of uninitialized variables. Setting up a local environment will be convenient for you to work on the project, but you know that in the end, your work should pass tests on the submit server. Make sure your work passes the test cases in the submit server.

1 Files

The first thing you should do is to check the files downloaded from ELMS. There are 8 files:

- `asm.c`

- `Makefile`

- `mips-small.c`

- `mips-small-pipe.c`

- `mips-small-pipe.h`

- `public_tests/*.mips`

- `public_tests/*.output`

Name	Format Type	Opcode	Func
LW	I-type	0x23	
SW	I-type	0x2B	
BEQZ	I-type	0x4	
ADDI	I-type	0x8	
ADD	R-type	0x0	0x20
SUB	R-type	0x0	0x22
SLL	R-type	0x0	0x4
SRL	R-type	0x0	0x6
AND	R-type	0x0	0x24
OR	R-type	0x0	0x25
HALT	J-type	0x3F	

Table 1: Instruction encodings for a reduced MIPS ISA.

The following is a short description on each file in the starter code:

- `asm.c` is an assembler for the reduced MIPS ISA which your simulator will implement (more about the assembler and the ISA later).

- `Makefile` is a unix make file which will produce two binaries `asm`, `sim`, and `sim-pipe` from the source files, `asm.c`, `mips-small.c`, and `mips-small-pipe.c`, respectively.

- `mips-small.c` is the C source file for a fully functional **unpipelined** simulator.

- `mips-small-pipe.c` is a C source code template that has some useful data structures and routines for a pipelined simulator. You can use this template to get started on the assignment.

- `mips-small-pipe.h` is a header file included by `mips-small-pipe.c`

- In the `public_tests` directory, you will find public test files. `*.mips` are the MIPS assembly programs and `*.output` files are the expected output by running the machine code generated from `*.mips` files. Note that you do NOT directly provide a `.mips` file to the simulator. You will first need to generate machine code using the assembler (`asm`) and the generated machine code is the input to the simulator.

2 A Scaled-Down MIPS ISA

You will be simulating the MIPS ISA from Hennessy & Patterson, with some key differences. First, instead of a 64-bit architecture, you will implement a 32-bit architecture. In other words, all registers and data paths are 32 bits wide, and all instructions will operate on 32-bit operands – i.e., *an integer variable (4 bytes) should be big enough to hold an instruction on the grace cluster.* Second, to keep your simulator simple, you will only be required to support a scaled-down version of the MIPS ISA consisting of 11 instructions. These instructions, along with their encoding, are given in Table 1. We will adhere to the MIPS instruction formats presented in Figure A.22 (Instruction layout for MIPS) of Hennessy & Patterson (5th edition), with *the exception that there is no "shamt" field in the R-type format, and instead the "func" field is 11-bits wide.* Finally, although all immediate values in branch and jump instructions are already left-shifted by 2, so your `BEQZ` instruction **SHOULD NOT** perform the left shift in the implementation. A simple implementation of MIPS in pages C-31 through C-33 and Figure C.23 would give you some hints on pipeline implementation, although the implementation is not identical to the requirements explained in this description.

Notice that all the instructions in Table 1 (except for one) exist in the normal MIPS ISA. These instructions behave exactly as described in the text (except they are 32-bit versions rather than 64-bit versions). The instruction we've added is the `HALT` instruction. As its name implies, when your simulator executes a HALT instruction, it should terminate the simulation. For more information about MIPS ISA, consult Section A.9 of Hennessy & Patterson.

3 asm: An Assembler for the Reduced MIPS ISA

We have provided an assembler, `asm.c`, so that you can assemble programs for your simulator. The `asm.c` file is fully functional, and you will not make any modifications to this file. Simply use the `Makefile` to make the binary `asm` from the `asm.c` source file – i.e., You would work on a Linux environment and type `make` to compile everything using the `Makefile`.

The format for assembly programs is very simple. A valid assembly program is an ASCII file in which each line of the file represents a single instruction, or a data constant. The format for a line of assembly code is:

```
label<tab>instruction<tab>field0<tab>field1<tab>field2<tab>comments
```

The leftmost field on a line is the `label` field which indicates a symbolic address. Valid labels contain a maximum of 6 characters and can consist of letters and numbers. The label is optional (the tab following the label field is not. After the optional label is a tab – i.e., – the tab character indicates if there is a label or not. Then follows the instruction field, where the instruction can be any of the assembly-language mnemonics listed in Table 1. After another tab comes a series of fields. All fields are given as **decimal** numbers. The number of fields depends on the instruction. The following describes the instructions and how they are specified in assembly code:

```
lw     rd    rs1    imm      Reg[rd] <- Mem[Reg[rs1] + imm]
sw     rd    rs1    imm      Reg[rd] -> Mem[Reg[rs1] + imm]
beqz   rd    rs1    imm      if (Reg[rs1] == 0) PC <- PC+4+imm
addi   rd    rs1    imm      Reg[rd] <- Reg[rs1] + imm
add    rd    rs1    rs2      Reg[rd] <- Reg[rs1] + Reg[rs2]
sub    rd    rs1    rs2      Reg[rd] <- Reg[rs1] - Reg[rs2]
sll    rd    rs1    rs2      Reg[rd] <- Reg[rs1] << Reg[rs2]
srl    rd    rs1    rs2      Reg[rd] <- Reg[rs1] >> Reg[rs2]
and    rd    rs1    rs2      Reg[rd] <- Reg[rs1] & Reg[rs2]
or     rd    rs1    rs2      Reg[rd] <- Reg[rs1] | Reg[rs2]
halt                         stop simulation
```

Note that in the case of the `beqz` instruction, PC-relative addressing is used (and again, your simulator should not perform the left-shift when computing the PC-relative branch target). For the `lw`, `sw`, and `beqz` instructions, the `imm` field can either be a decimal value, or a label can be used. In the case of a label, the assembler performs a different action depending on whether the instruction is a `lw` / `sw` instruction, or a `beqz` instruction. For `lw` and `sw` instructions, the assembler inserts the absolute address corresponding to the label. For `beqz` instructions, the assembler computes a PC-relative offset with respect to the label. After the last field is another tab, then any comments. The comments end at the end of the line.

In addition to instructions, lines of assembly code can also include directives for the assembler. The only directive we will use is `.fill`. The `.fill` directive tells the assembler to put a number into the place where the instruction would normally be stored. The `.fill directive` uses one field, which can be either a numeric value or a symbolic address. For example, in the following code, ".fill 32" puts the value 32 where the instruction would normally be stored. In the following example, ".fill start" will store the value 8, because the label "start" refers to address 8 (remember that the MIPS architecture uses byte addressable memory).

<u>An example in the next page</u>

```
        addi    1    0    5        load reg1 with 5
        addi    2    0    -1       load reg2 with -1
start   add     1    1    2        decrement reg1
        lw      3    0    var1     loads reg3 with value stored in var1
        addi    3    3    -1       decrement reg3
        sw      3    0    var1     put reg3 back (thus var1 is decremented)
        beqz    0    1    done     goto done when reg1==0
        beqz    0    0    start    back to start
        add     0    0    0
done    halt
        .fill   start                     will contain start address (8)
var1    .fill   32                        Declare a variable, initialized to 32
```

Try taking the above example and running the assembler on it. Enter the above assembly code into a file called `ex.mips`. Then type `asm ex.mips ex.out`. The assembler will generate a file `ex.out` which should contain:

```
(address 0x0):    20010005
(address 0x4):    2002ffff
(address 0x8):    00220820
(address 0xc):    8c03002c
(address 0x10):   2063ffff
(address 0x14):   ac03002c
(address 0x18):   10200008
(address 0x1c):   1000ffe8
(address 0x20):   00000020
(address 0x24):   fc000000
(address 0x28):   00000008
(address 0x2c):   00000020
```

The assembler assumes that all programs will be loaded into memory beginning at address 0x0.

4 Simulator without Pipelining

As described above, we have provided you with a simulator that simulates the same ISA, but which is not cycle accurate because it does not account for pipelining. The source code for this simulator is in `mips-small.c`. You can use this simulator to run assembly language programs prior to getting the pipelined simulator working.

Build the unpipelined simulator by typing `make sim` in the directory where you've copied the files we have provided. This will produce an executable `sim` which you can run. Try running this simulator on the program we have provided, `publicMult.mips`. This program multiplies two numbers, specified in memory at the labels `mcand` and `mplier`, and stores the result at the label `answer`. Don't forget to assemble the program before running the simulator. The sequence of commands you should give are:

```
make
asm public_tests/publicMult.mips public_tests/publicMult.machine
sim public_tests/publicMult.machine
```

Note that `make sim` will produce the binary code for `mips-small.c` only. If you want to generate binary for other source code, check the targets in `Makefile` – `sim` is a target in the file.

When you run the simulator, a ton of messages will spew onto your terminal. The simulator we have provided dumps the state of the machine (including the contents of registers and memory) at every simulated cycle. Note that `mips-small.c` implements an single-cycle, unpipelined machine. So, each instruction completes in a cycle, which is *long* enough to process everything needed from fetching an instruction through updating the register file.

4

You can examine this output to see what the simulator is doing. To put these messages in a file so that you can actually read them, redirect the output of the simulator to a file. – e.g., on a C shell, by executing the following command:

```
sim public_tests/publicMult.machine >! publicMult.myout
```

The syntax for input-output redirection can be different depending on the shell you use. You can identify the shell type using the following command:

```
env | grep SHELL
```

The use of I/O redirection will be very useful (and really important) as you work on this project, because you can see differences between an output file produced and the expected output file (e.g., using the `diff` command).

Your *pipelined* simulator (`mips-small-pipe`) will spew similar messages, except it will also include the state of pipeline registers. It is recommended to understand the unpipelined simulator code before you move on to building the pipelined simulator.

5 Assignment

Your assignment is to build a cycle-accurate simulator that accounts for pipelining. We have provided starter code in `mips-small-pipe.c` and `mips-small-pipe.h`. This code serves two purposes: it will make your life a bit easier since you won't have to write the whole simulator from scratch, and it is meant to enforce some coding disciplines that ensure you actually simulate the details of the processor pipeline.

The following sections describe the assignment in more detail.

5.1 Simulate Pipelining

You will be simulating the basic MIPS pipeline depicted in Figure C.22 of Hennessy & Patterson. We have provided pipeline register data structures specified in the file, `mips-small-pipe.h`. **To enforce an implementation of your simulator that simulates pipelining, you are required to use these structures in their unmodified form.** In other words, each cycle simulated by your simulator must produce the correct values for all the fields in the pipeline register data structures.

```
typedef struct IFIDStruct {
  int instr;
  int pcPlus1;
} IFID_t;

typedef struct IDEXStruct {
  int instr;
  int pcPlus1;
  int readRegA;
  int readRegB;
  int offset;
} IDEX_t;

typedef struct EXMEMStruct {
  int instr;
  int aluResult;
  int readRegB;
} EXMEM_t;
```

```
typedef struct MEMWBStruct {
  int instr;
  int writeData;
} MEMWB_t;

typedef struct WBENDStruct {
  int instr;
  int writeData;
} WBEND_t;

typedef struct stateStruct {
  int pc;
  int instrMem[MAXMEMORY];
  int dataMem[MAXMEMORY];
  int reg[NUMREGS];
  int numMemory;
  IFID_t IFID;
  IDEX_t IDEX;
  EXMEM_t EXMEM;
  MEMWB_t MEMWB;
  WBEND_t WBEND;
  int cycles; /* number of cycles run so far */
} state_t, *Pstate;
```

There are three small modifications in the above pipeline register data structures as compared with Figure C.22 in Hennessy & Patterson. First, we don't latch the "branch taken" signal in the EXMEM register. Instead, when your simulator computes the direction of the branch in the execute (EX) stage, you can directly modify the PC register from the execute (EX) stage. Second, the MUX in the WB stage has been moved to the MEM stage. Therefore, the MEMWB pipeline registers only latch two values, `instr` and `writeData`, instead of three values as indicated by Figure C.22 in the textbook. Finally, we have added an extra pipeline register after the WB stage. Instead of communicating internally through the register file, your simulator will always forward from the pipeline registers. For instance, in Figure C.7 of the textbook, there is a communication path from the DADD instruction to the OR instruction internally through the register file (this is achieved by writing the register file on the first half of the clock cycle, and reading the register file on the second half of the clock cycle). In your simulator, this register communication will be replaced by forwarding between the WBEND pipeline registers to the execute stage of the OR instruction in cycle CC6.

As mentioned above, your simulator must produce all the values for the pipeline register data structures on a cycle-by-cycle basis. **In addition, your simulator must also use the following simulator loop code** (which is in `mips-small-pipe.c`):

```
while (1) {

    printState(state);

    /* copy everything so all we have to do is make changes.
       (this is primarily for the memory and reg arrays) */
    memcpy(&new, state, sizeof(state_t));

    new.cycles++;

    /* -------------------- IF stage -------------------- */
```

6

```
    /* -------------------- ID stage -------------------- */

    /* -------------------- EX stage -------------------- */

    /* -------------------- MEM stage -------------------- */

    /* -------------------- WB stage -------------------- */

    /* transfer new state into current state */
    memcpy(state, &new, sizeof(state_t));
}
```

You will insert the necessary code for the stages in the while loop (between two `memcpy` statements). **Be careful not to modify the output format of printState.**[1] Within each pipeline stage, the code will compute the new machine state – i.e., memory state and pipeline register data structures. The new state should be stored in the structure named "`new`") as a function of the old values (stored in the structure named "`state`").

After the code in all the stages have computed their values, the `memcpy` at the end of the simulator loop copies the values from the `new` structure into the `state` structure. This corresponds to the rising edge of the clock latching values into the pipeline registers for the next clock cycle.

It is important to understand that the stages simulate independent hardware units, although we put code for the stages in-order in the loop. Assume that your code for each stage is run at the same time, although they are run top-down in reality. Because the stages must be independent, avoid using local variables declared for an earlier stage.

Also note that new machine state (`new`) is determined based on the current machine state (`state`). That is, the variable `new` should appear on the left-hand-side of any assignment statements, and `state` should appear on the right-hand-side.

> ***We will manually check your code and penalize:***

- local variable(s) are declared for a stage and used for other stage(s)

- if the variable `state` appears in the left-hand-side, or if `new` appears in the right-hand-side of any assignment statements

5.2 Dealing with Data Hazards

Your simulator should deal with data hazards by **bypassing (or forwarding)** values from the pipeline register data structures. You can implement such data forwarding by checking for data dependences in the execute stage. For every instruction that enters the execute stage, you will have to check against the results produced by the previous 3 instructions (whose results can be found in the EXMEM, MEMWB, and WBEND pipeline registers of the old machine state). If there are no data dependences, then the execute stage should take its operands from the `readRegA` and `readRegB` in the IDEX pipeline register. If there are data dependences, then the execute stage should instead copy the results from the pipeline register of the dependent instruction.

Most data hazards can be resolved using bypassing. However, there is one case that requires stalling the pipeline. As seen in Figure C.9 in the textbook, if an instruction uses a source register that is the destination register of an immediately preceding `load` instruction, then the pipeline must be stalled for a single cycle. This can be implemented by checking for such a dependence in the decode stage and inserting a bubble (NOP instruction - see Section 5.4.2) instead of allowing the instruction to proceed. One special case is when a load from a location is immediately followed by a store (Figure C.8 in the textbook). The depending store can be bypassed by forwarding from the MEM/WB latch (loaded value) to the data memory write port of the memory stage (store instruction). However, for your implementation, you will not implement this; simply stall the pipeline one clock cycle for ALL

---

[1]You are allowed to write helper functions in the C file if needed.

load instructions that are followed by a depending instruction, even if the depending instruction is a store.

5.3   Dealing with Control Hazards

The pipeline you will be simulating will have a branch delay of 2 cycles since the branch direction is not resolved until the end of the execute stage. Instead of stalling the pipeline for 2 cycles after every fetched branch, your simulator should implement a "hybrid" static branch prediction scheme. In this scheme, you should predict taken for backward branches (i.e., with a negative offset), and predict not-taken for forward branches. To implement this hybrid scheme, **test the immediate value in the fetch stage of every fetched branch instruction.** If the immediate value is positive, then fetch the next consecutive instruction in the next cycle. If the immediate value is negative, then fetch the branch target (computed as a PC-relative address using the offset) in the next cycle.[2]

Make sure that your simulator validates the branch direction computed in the execute stage against the predicted direction. If the prediction was correct, no action should be taken. If the prediction was incorrect, then the simulator must squash the two instructions in the pipeline, which at this point will be in the IFID and IDEX pipeline registers. This squashing can be accomplished by replacing the incorrectly fetched instructions with NOP instructions and clearing all the other state in the IFID and IDEX pipeline registers.

5.4   Miscellaneous

5.4.1   Memory

As indicated by the `state_t` structure described in Section 5.1, there are two arrays that implement memory, an instruction memory array `instrMem` and a data memory array `dataMem`. This simulates the separate instruction and memory structures in the pipeline in Figure C.22 in the textbook. The code we have provided initializes both of these arrays to the same values read in from the assembly code file. In your simulator, you should read instructions from the `instrMem` array in the fetch stage, and write (read) values to (from) the `dataMem` array in the memory stage.

In the simulator code we have provided for you, the amount of memory simulated is determined by the size of the assembly program file that you read into the simulator. The `numMemory` field of the `state_t` structure is set to this size during initialization and is fixed for the duration of the simulation.[3] This means that you must allocate all the memory needed to run your program statically in the assembly code file. The maximum size of the assembly program allowed is set by the `MAXMEMORY` constant declared in `mips-small-pipe.h` and is currently 16K words (64K bytes).

5.4.2   NOP Instruction

The MIPS ISA does not include an explicit NOP (null operation) instruction, and neither does our reduced MIPS ISA. However, the NOP instruction is needed to initialize all the instruction fields of the pipeline register data structures, and to replace or "squash" those instructions in the pipeline following a branch misprediction. In your simulator, you will use the "`add 0 0 0`" instruction as a NOP instruction (since this instruction does not effect any state). We have already declared this as `NOPINSTRUCTION` and used it in the initialization code.

5.4.3   Halting

When your simulator encounters the `HALT` instruction, it should end the simulation. Notice, however, that it is incorrect to halt as soon as the HALT instruction is fetched since this code may be speculative (dependent on a yet to be verified branch). If the branch was mis-predicted, then you definitely do not want to halt. Another issue is that any `SW` instructions which are farther along

---

[2]So, for a correct taken/not-taken prediction, no cycle will be wasted.
[3]See the main function

the pipeline than the `HALT` instruction must be allowed to complete before the machine is stopped. To solve these problems, halt the machine when the `HALT` instruction reaches the MEMWB pipeline register. This ensures that all previously executed instructions have completed.

Note that the emulator will continue to fetch instructions into the pipeline, even after the HALT instruction. They are not real instructions but simply values stored in `instrMem`. As described in Section 5.4.1, the code we have provided initializes both of these arrays to the save values read in from the assembly code file. If the `HALT` instruction is the last line in the assembly code file, the emulator may read 0 from `instrMem` (the `int` array in the emulator implementation). This is okay because the machine will halt before the instructions read after the `HALT` instruction make any changes to the system.

5.5 Grading

Your work will be tested against the tests in the submit server. Do NOT modify the simulator loop structure and the `printState` function that in the distributed file. The printed output will be compared by the submit server.

Each time you call `printState`, it should print in the format that we have provided, and you should only call `printState` once every iteration of the simulator loop (you are free to use `printState` as much as you want when debugging, but take all the extra `printState` calls out of your code before handing in your assignment).

The public tests and the expected output files are provided. You can compare your program output with the expected output using input-output redirection and the `diff` tool. For example, if the assembly code generated from `simple.mips` is executed and if the printed output is saved to `simple.out`, There should be no differences between `simple.out` and the expected output (say the file name `simple.output`), when you run `diff simple.output simple.out`.

Submit your implementation in two parts (Phase 1 and Phase 2). In the Phase 1 submission (40% of the project points), you will complete the basic structure of the pipeline simulator that can process ADD, ADDI, LW, and `HALT` instruction, without implementing forwarding and branching. You must implement all features in the Phase 2 submission. (60% of the project points).

5.6 Academic Honesty

Do not allow any other student to see any of your code. You may however discuss the assignment in general terms, with other students, only for the clarification purpose. The line between clarification purpose discussion and academic dishonesty is thin. Discussing too much implementation details will make your work very similar to another student's work and it will be considered as excessive collaboration. If copying or excessive collaboration is detected in your submissions, the matter will be referred to the Office of Student Conduct.

6 Submission

You will submit `mips-small-pipe.h` and `mips-small-pipe.c` to the submit server. **Note that you will need to incrementally implement and submit two versions. Read the Grading section for more details.**